

## High level design

The client, *workerctl*, is a command line interface that talks to the worker server over HTTPS, meaning there is no persistent connection between them. After issuing a command, for example *workerctl start "ls -l"*, the client sends an async POST request to the server.

On the server, after authenticating the user request, a goroutine gets created to handle the request, and, since commands may take a long time to execute, the server immediately sends a response to the client that contains the jobID and status "created".

To get status/output of the process, the client will send another request with the jobID returned from the previous command.

All job goroutines will have a timeout of 60 seconds to prevent them from running too long.

Server representation of a job:

```
type job struct {  
    mu      sync.RWMutex  
    id      string  
    userid  string  
    command string  
    status  string  
    output  []string  
}
```

## Security

**Authentication** will be handled by a pre-generated token saved in the repo. In the real world we would want to use an authentication service/package and a database to store tokens after login. Each API request will be authenticated first, before performing the request.

**Authorization** is the same for all registered users, they all can start/stop/check status of their own processes only.

All interaction between the client and the server are over **HTTPS**, where the client validates the server's credentials. The server doesn't care about the client's identity.

## API

All API requests will be POST. The difference between requests will be in the body.

**Request and Response Parameters:**

action - ["start", "stop", "status"]

status - ["created", "running", "finished", "stopped", "error"]  
token - authentication and authorization string  
output - process output from stdout and stderr  
error - API error message

### **Start a job**

**request:** {"action": "start", "token": <token>, "command": <command>}  
**success response:** {"jobID": <jobID>, "status": "created"}  
**error response:** {"error": <errorMessage>}

### **Stop a job**

**request:** {"action": "stop", "token": <token>, "jobID": <jobID>}  
**success response:** {"jobID": <jobID>, "status": "stopped"}  
**error response:** {"error": <errorMessage>}

A job can only be stopped from the "created" or "running" states. if the job is in the "finished" or "error" states, it remains unchanged.

### **Job status**

**request:** {"action": "status", "token": <token>, "jobID": <jobID>}  
**success response:** {"jobID": <jobID>, "status": <status>, "output": <output>}  
**error response:** {"error": <errorMessage>}

## **Client**

All requests sent by the client are asynchronous, allowing for submitting a batch of commands. It will be the responsibility of the user to track requests and responses.

### **Start a job**

```
workerctl start "ls -l"  
jobID: 7asd79as87d9asd86d  
status: created
```

### **Stop a job**

```
workerctl stop <jobID>  
jobID: 7asd79as87d9asd86d  
status: stopped
```

### **Job status**

```
workerctl status <jobID>  
jobID: 7asd79as87d9asd86d  
status: finished  
output:  
drwx-----@  3 user  staff    96 28 Jan 20:59 Applications  
drwx-----@ 27 user  staff   864 14 Apr 11:57 Desktop
```

## Scope and Trade-offs

### **To write as little code as possible, I will:**

- hardcode a secret token for authentication
- write one or two tests for authentication and networking
- not isolate user environments on the server
- not create a configuration file
- not make it scalable or highly available

### **Future considerations:**

For a real-world product, we would want each user to have their own isolated environment for running Linux commands. This could be achieved by using docker containers. After a user authenticates with the server for the first time, a docker container would be spun up and assigned to the user. After some time of inactivity, the container would be stopped, preserving the state for future requests. To prevent resource hogging we would limit resources assigned to each container.

For a more real-time experience, we would want the server to send updates to the client whenever there is new output from a process, and for that we could use gRPC.